



# Extending the project to perform both tool calls as well as analyzing the responses based on the tool outputs

Short answer first: **yes**, this makes mathematical sense and **yes**, it's fully compatible with SkyRL.

- In RL terms, you're defining a **dense per-turn reward** for tool usage + analysis **and** a **terminal reward** that scores the **final text answer** against a **reference answer** produced by your synthetic data generator. Maximizing expected return under GRPO/PPO will push the policy toward correct tools/arguments *and* high-quality final summaries.
- SkyRL is built for this: multi-turn `BaseTextEnv` tasks, tool calls via ToolGroups, final answers (e.g., the Search example uses `<answer>...</answer>` and gives reward for correctness), and dataset entries that keep only a seed prompt while the environment computes turn-level rewards. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

Below is a **complete recipe**:

---

## 0) What “mathematically makes sense” here

Let an episode be  $(\tau = (s_0, a_0, r_0, \dots, s_T, a_T, r_T))$  where:

- Steps  $(1..K)$  are **tool turns**; step  $(K+1)$  is the **final text answer**.
- Reward  $(r_t = r_t^{\{\text{tools}\}})$  for tool turns (shaped by your rubric), and  $(r_{K+1} = \alpha \cdot r^{\{\text{final}\}})$  where  $(r^{\{\text{final}\}} \in [0, 1])$  comes from **LLM-as-a-Judge** (LAJ) comparing the **policy's final answer** vs the **reference** (and/or a set of “verifiable facts” extracted during data generation).  $(\alpha)$  balances tool learning vs answer quality.

GRPO/PPO maximizes  $(\mathbb{E}_{\pi}[\sum_t r_t])$  using advantages  $(A_t)$  (e.g., GAE). Since LAJ outputs are **numeric rewards**, they're just another term in the return; no special math is needed. (This is analogous to SkyRL's Search example, which assigns a terminal 0/1 for correct/incorrect final text. We're just using a **graded** judge instead of binary

correctness.) ([skyr1.readthedocs.io](https://skyr1.readthedocs.io))

---

## 1) Data model you should write to `train_llm.json`

Add **two** new sections to your `reward_spec.ground_truth` :

- `final_reference` : the **reference final answer** your generator produces *by actually running the plan over MCP tools and analyzing the outputs*. Store both **text** and **facts** (structured), plus optional citations.
- `judge_rubric` : a **structured rubric** your environment will pass to LAJ to score the policy's final text against the reference.

## **1.1 Minimal schema extension (drop-in)**

```

{
  "data_source": "synthetic/llm",
  "env_class": "MCPToolEnv",
  "prompt": [ { "role": "system", "content": "..." }, { "role": "user", "content": "..." } ],
  "reward_spec": {
    "method": "rule",
    "ground_truth": {
      "task_id": "string",
      "complexity": "simple|moderate|complex",
      "max_turns": 10,
      "limits": { "max_servers": 3, "max_tools": 10 },
      "tool_sequence": [ { "step": 1, "server": "...", "tool": "...", "params": { } }, "..." ],
      "analysis_rubric": {
        "steps": [ { "step": 1, "extract": [], "compute": [], "select": [], "accept_if":
          "final_answer_requirements": {
            "format": "text|markdown|json",
            "must_include": ["list","of","keys-or-names"],
            "grounded_from": ["state_keys_to_check"],
            "quality_criteria": ["no hallucinations","concise", "..."]
          }
        }
      },
    },
    "final_reference": {
      "answer_text": "the reference final summary, generated by your agent after running
      "facts": { "top3": ["NVDA","AMD","META"], "neg_titles": ["...", "..."] },
      "citations": { "top3": [3], "neg_titles": [8,9] } // step indexes that support f
    },
    "judge_rubric": {
      "weights": { "coverage": 0.3, "grounding": 0.4, "clarity": 0.2, "safety": 0.1 },
      "schema": {
        "type": "object",
        "properties": {
          "coverage": { "type": "number", "minimum": 0, "maximum": 1 },
          "grounding": { "type": "number", "minimum": 0, "maximum": 1 },
          "clarity": { "type": "number", "minimum": 0, "maximum": 1 },
          "safety": { "type": "number", "minimum": 0, "maximum": 1 },
          "total": { "type": "number", "minimum": 0, "maximum": 1 }
        },
        "required": ["coverage", "grounding", "clarity", "safety", "total"]
      },
    },
  },

```

```

        "target_length_range": [40, 140]
    }
}
},
"extra_info": { "scenario": {"scenario": "...", "turns": 10} }
}

```

- Keep the **prompt compact** (system + first user) per SkyRL's Dataset Preparation guidance. The env drives multi-turn reward logic. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- `final_reference` is **not** shown to the policy; it's used only for reward.

## 2) Changes to your generator

### `src/dataset/llm/generate_with_llm.py`

You'll add **two phases** after planning:

1. **Execute the plan** over MCP tools, step-by-step, applying your **analysis\_rubric.steps** (extract/compute/select/accept\_if/next\_args\_from) to build a **named state** and capture **per-step result summaries**.
2. **Compose the reference final text** using the named state (either via deterministic templating or by calling an LLM with a short, grounded prompt), and write it as `final_reference.answer_text` (+ `facts`, `citations`).

Tip: keep the **DSL** small and safe (no `eval`), e.g., functions like `last`, `prev`, `pct_change_last_day`, `topk`, `argmax`, `unique`, `regex_extract_all`, `concat`, `head`. This mirrors SkyRL's practice of keeping env logic deterministic and light. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

## **2.1 Minimal, PR-ready patch (conceptual diff)**

```

--- a/src/dataset/llm/generate_with_llm.py
+++ b/src/dataset/llm/generate_with_llm.py
@@
-TASK_SCHEMA = {...} # your existing schema
+TASK_SCHEMA = {
+    "name": "rl_task",
+    "schema": {
+        "type": "object",
+        "properties": {
+            "task_id": {"type": "string"},
+            "complexity": {"type": "string", "enum": ["simple", "moderate", "complex"]},
+            "user_prompt": {"type": "string"},
+            "max_turns": {"type": "integer", "minimum": 3, "maximum": 16},
+            "tools_available": {"type": "array", "items": {"type": "string"}},
+            "limits": {"type": "object"},
+            "tool_sequence": {
+                "type": "array", "minItems": 2, "maxItems": 12,
+                "items": {
+                    "type": "object",
+                    "properties": {
+                        "step": {"type": "integer", "minimum": 1},
+                        "server": {"type": "string"},
+                        "tool": {"type": "string"},
+                        "params": {"type": "object"},
+                        "analysis_requirements": {
+                            "type": "object",
+                            "properties": {
+                                "extract": {"type": "array", "items": {"type": "string"}},
+                                "compute": {"type": "array", "items": {"type": "string"}},
+                                "select": {"type": "array", "items": {"type": "string"}},
+                                "accept_if": {"type": "array", "items": {"type": "string"}},
+                                "next_args_from": {"type": "string"}
+                            },
+                            "required": ["next_args_from"],
+                            "additionalProperties": true
+                        }
+                    },
+                    "required": ["step", "server", "tool", "params", "analysis_requirements"],
+                    "additionalProperties": true
+                }
+            }
+        }
+    }

```

```

+     }
+ },
+     "final_answer_requirements": {
+         "type": "object",
+         "properties": {
+             "format": {"type": "string"},
+             "must_include": {"type": "array", "items": {"type": "string"}},
+             "grounded_from": {"type": "array", "items": {"type": "string"}},
+             "quality_criteria": {"type": "array", "items": {"type": "string"}}
+         },
+         "required": ["format", "must_include", "grounded_from"]
+     },
+     "judge_rubric": {
+         "type": "object",
+         "properties": {
+             "weights": {"type": "object"},
+             "schema": {"type": "object"},
+             "target_length_range": {"type": "array", "items": {"type": "integer"}}
+         },
+         "required": ["weights", "schema"]
+     }
+ },
+     "required": ["task_id", "complexity", "user_prompt", "max_turns", "tool_sequence", "final_
+     "additionalProperties": false
+ }
+}

```

@@

```

async def _one_task(...):
    # 1) Ask LLM for plan (as you do today)
    task = await _call_llm_with_schema(...)
-    return task
+    # 2) Verify & repair chaining (next_args_from, step ranges)
+    task = _verify_and_repair(task)
+    # 3) Execute plan over MCP tools to build named state and per-step summaries
+    exec_out = await simulate_plan_and_collect(task, tool_manager)
+    # 4) Compose reference final answer (LLM or template), grounded in exec_out.state
+    final_ref = await compose_reference_answer(task, exec_out)
+    task["_exec_out"] = exec_out.to_dict()    # optional: keep light summaries, hashes
+    task["_final_reference"] = final_ref
+    return task

```



@@

```
-def to_skyrl_sample(task, system_prompt):  
-    ground_truth = {...}  
+def to_skyrl_sample(task, system_prompt):  
+    ground_truth = {  
+        "task_id": task["task_id"],  
+        "complexity": task["complexity"],  
+        "max_turns": task["max_turns"],  
+        "limits": task.get("limits", {}),  
+        "tool_sequence": task["tool_sequence"],  
+        "analysis_rubric": {  
+            "steps": [{ "step": s["step"], **s["analysis_requirements"] } for s in task["tool_sequence"]],  
+            "final_answer_requirements": task["final_answer_requirements"]  
+        },  
+        "final_reference": task["_final_reference"],  
+        "judge_rubric": task["judge_rubric"]  
+    }  
    return {  
        "data_source": "synthetic/llm",  
        "env_class": "MCPToolEnv",  
        "prompt": [  
            {"role": "system", "content": system_prompt},  
            {"role": "user", "content": task["user_prompt"]}  
        ],  
-        "reward_spec": {"method": "rule", "ground_truth": ground_truth},  
+        "reward_spec": {"method": "rule", "ground_truth": ground_truth},  
        "extra_info": {"version": "lh-v2"}  
    }
```

## **2.2 Execution harness (new helpers)**

```

async def simulate_plan_and_collect(task: dict, tm) -> "ExecOut":
    """
    Runs tool_sequence with placeholder resolution and applies analysis_rubric.steps
    to produce named state and lightweight per-step summaries.
    """
    plan = task["tool_sequence"]
    steps_rubric = [s["analysis_requirements"] for s in plan]
    state = {} # named values introduced by extract/compute/select
    per_step = []
    for idx, step in enumerate(plan, 1):
        args = resolve_placeholders(step["params"], state) # ${var} replacement
        result = await tm.execute_tool(f"{step['server']}.{step['tool']}", args, timeout=)
        summary = summarize_tool_result(result) # small, e.g., top keys, co
        # apply rubric: extract/compute/select/accept_if
        updates, checks = apply_analysis(steps_rubric[idx-1], result, state)
        state.update(updates)
        per_step.append({"step": idx, "args": args, "summary": summary, "checks": checks})
    return ExecOut(state=state, steps=per_step)

def apply_analysis(ar: dict, result: dict, state: dict) -> tuple[dict, dict]:
    """Implements a tiny, safe DSL (no eval) for extract/compute/select/accept_if."""
    updates = {}
    checks = {"accept_pass": True, "missing": []}
    # extract
    for name in ar.get("extract", []):
        val, ok = safe_extract(name, result) # support "close[]", "articles[][title]"
        if ok: updates[name.split("[")[0]] = val
        else: checks["missing"].append(name); checks["accept_pass"] = False
    # compute/select (whitelisted ops like last, prev, pct_change_last_day, topk, argmax,
    for expr in ar.get("compute", []): updates.update(safe_compute(expr, state))
    for expr in ar.get("select", []): updates.update(safe_compute(expr, state))
    # accept_if
    for cond in ar.get("accept_if", []):
        if not safe_check(cond, state): checks["accept_pass"] = False
    return updates, checks

async def compose_reference_answer(task: dict, exec_out: "ExecOut") -> dict:
    """Produce final_reference {answer_text, facts, citations} from state."""
    far = task["final_answer_requirements"]

```

```
# build facts deterministically from state:
facts = build_facts(exec_out.state, far.get("grounded_from", []))
# use a small LLM prompt or a template to compose the text
answer_text = await small_llm_compose(far, exec_out.state, facts)
# map which steps support which facts
citations = infer_citations(facts, exec_out.steps)
return {"answer_text": answer_text, "facts": facts, "citations": citations}
```

- **Why do this at data-gen time?** So your **reference** is grounded in **real tool results**, not a hallucinated answer. At training, the env will compare the **policy's** final text to this reference using LAJ and/or heuristics.

## 3) Environment changes (how the final answer is rewarded)

Your `MCPToolEnv(BaseTextEnv)` should accept either:

- a **tool call** ( `{"tool": "server.tool", "arguments": {...}}` ) or
- a **final answer** ( `{"final_answer": "...text..."}` ), similar to the **Search example**, which uses `<answer>...</answer>` with stop strings to mark the last turn. ([skyr1.readthedocs.io](https://skyr1.readthedocs.io))

### 3.1 Final answer scoring in the env

**Heuristic (fast):**

- **Coverage:** does text include all names in `analysis_rubric.final_answer_requirements.must_include` ?
- **Grounding:** is text consistent with `final_reference.facts` (e.g., set overlap equals 1.0), and **does not** contradict them?
- **Clarity:** target length range; simple readability checks.
- **Safety:** blacklist checks.

**LLM-as-a-Judge (LAJ):**

- Judge sees a **compact state summary**, the **policy's final text**, and the **reference** (text +

facts).

- Use **Structured Outputs** with your `judge_rubric.schema` to force **numbers-only**.

### Reward:

- $(r_{\text{final}} = \lambda_{\text{heur}} r_{\text{heur}} + \lambda_{\text{la}} r_{\text{la}})$  (both in  $[0,1]$ ).
- Total episode return is sum of shaped tool rewards +  $(\alpha \cdot r_{\text{final}})$ .

This mirrors SkyRL's pattern where the final text inside `<answer>` is what gets rewarded (binary in Search; graded here). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

---

#### **4) Two sample dataset items (with final\_reference & judge\_rubric )**

**4.1 NASDAQ-100 news triage (truncated for space; drop  
into your JSON array)**

```

{
  "data_source": "synthetic/llm",
  "env_class": "MCPToolEnv",
  "prompt": [
    { "role": "system", "content": "You have tools DuckDuckGo, yahoo_finance, python_execution" },
    { "role": "user", "content": "Find top-3 gainers in NASDAQ-100 today, get 5 news headlines" },
  ],
  "reward_spec": {
    "method": "rule",
    "ground_truth": {
      "task_id": "nasdaq100_neg_digest_v2",
      "complexity": "complex",
      "max_turns": 12,
      "limits": { "max_servers": 3, "max_tools": 10 },
      "tool_sequence": [
        { "step": 1, "server": "DuckDuckGo", "tool": "search", "params": { "query": "NASDAQ-100 top gainers" } },
        { "step": 2, "server": "DuckDuckGo", "tool": "fetch_content", "params": { "url": "https://www.duckduckgo.com/search?q=NASDAQ-100+top+gainers" } },
        { "step": 3, "server": "yahoo_finance", "tool": "get_yfinance_price_history", "params": { "tickers": "top3" } },
        { "step": 4, "server": "python_execution", "tool": "python_execution", "params": { "code": "import json; print(json.dumps(top3))" } },
        { "step": 5, "server": "DuckDuckGo", "tool": "search", "params": { "query": "${top3[0]} news" } },
        { "step": 6, "server": "DuckDuckGo", "tool": "search", "params": { "query": "${top3[1]} news" } },
        { "step": 7, "server": "DuckDuckGo", "tool": "search", "params": { "query": "${top3[2]} news" } },
        { "step": 8, "server": "python_execution", "tool": "python_execution", "params": { "code": "import json; print(json.dumps(news_titles))" } },
        { "step": 9, "server": "python_execution", "tool": "python_execution", "params": { "code": "import json; print(json.dumps(neg_titles))" } },
        { "step": 10, "server": "slack", "tool": "send_slack_message", "params": { "channel": "#nasdaq100", "text": "Top 3 gainers and news headlines" } },
      ],
    },
    "analysis_rubric": {
      "steps": [
        { "step": 1, "extract": ["tickers_url"], "compute": [], "select": [], "accept_if": [] },
        { "step": 2, "extract": ["content"], "compute": ["tickers = regex_extract_all('https://www\\.duckduckgo\\.com/search?q=NASDAQ-100+top+gainers')"], "select": [], "accept_if": [] },
        { "step": 3, "extract": ["price_json"], "compute": ["pct = pct_change_last_day(pct_change_last_day(json.loads(price_json)))"], "select": [], "accept_if": [] },
        { "step": 4, "extract": ["top3[]"], "compute": [], "select": [], "accept_if": ["top3[0] != top3[1] != top3[2]"] },
        { "step": 5, "extract": ["articles0[][title]"], "compute": ["news_titles0 = title_sentiment_map(title_sentiment_map(json.loads(articles0[0]['content'])))"], "select": [], "accept_if": [] },
        { "step": 6, "extract": ["articles1[][title]"], "compute": ["news_titles1 = title_sentiment_map(title_sentiment_map(json.loads(articles1[0]['content'])))"], "select": [], "accept_if": [] },
        { "step": 7, "extract": ["articles2[][title]"], "compute": ["news_titles2 = title_sentiment_map(title_sentiment_map(json.loads(articles2[0]['content'])))"], "select": [], "accept_if": [] },
        { "step": 8, "extract": ["title_sentiment_map{title->score}"], "compute": ["news_titles = news_titles0 + news_titles1 + news_titles2"], "select": [], "accept_if": [] },
        { "step": 9, "extract": ["neg_titles[]"], "compute": [], "select": ["neg_titles"], "accept_if": [] },
        { "step": 10, "extract": [], "compute": [], "select": [], "accept_if": [], "next_step": 11 },
      ],
    },
  },
}

```

```

    "final_answer_requirements": {
      "format": "markdown",
      "must_include": ["top3", "neg_titles"],
      "grounded_from": ["top3", "title_sentiment_map"],
      "quality_criteria": ["relevant headlines", "no hallucinated tickers", "concise"]
    },
  },
  "final_reference": {
    "answer_text": "Top-3 NASDAQ-100 gainers today: NVDA, AMD, META. Notable negative
    "facts": { "top3": ["NVDA","AMD","META"], "neg_titles": ["...", "...", "..."] },
    "citations": { "top3": [4], "neg_titles": [8,9] }
  },
  "judge_rubric": {
    "weights": { "coverage": 0.35, "grounding": 0.4, "clarity": 0.15, "safety": 0.10 },
    "schema": {
      "type": "object",
      "properties": {
        "coverage": {"type": "number", "minimum": 0, "maximum": 1},
        "grounding": {"type": "number", "minimum": 0, "maximum": 1},
        "clarity": {"type": "number", "minimum": 0, "maximum": 1},
        "safety": {"type": "number", "minimum": 0, "maximum": 1},
        "total": {"type": "number", "minimum": 0, "maximum": 1}
      },
      "required": ["coverage", "grounding", "clarity", "safety", "total"]
    },
    "target_length_range": [40, 140]
  }
}
}
}
}

```



## **4.2 S3 error histogram (final answer includes decision & justification)**

```

{
  "data_source": "synthetic/llm",
  "env_class": "MCPToolEnv",
  "prompt": [
    { "role": "system", "content": "You have tools aws, python_execution, slack, jira. Em"},
    { "role": "user", "content": "Find buckets >10 GB, compute today's error %, post a hi"},
  ],
  "reward_spec": {
    "method": "rule",
    "ground_truth": {
      "task_id": "aws_error_histogram_v2",
      "complexity": "complex",
      "max_turns": 12,
      "limits": { "max_servers": 4, "max_tools": 10 },
      "tool_sequence": [
        { "step": 1, "server": "aws", "tool": "aws_s3_list_buckets", "params": {} },
        { "step": 2, "server": "python_execution", "tool": "python_execution", "params": {} },
        { "step": 3, "server": "aws", "tool": "aws_s3_list_objects", "params": { "bucket": "" },
        { "step": 4, "server": "python_execution", "tool": "python_execution", "params": {} },
        { "step": 5, "server": "aws", "tool": "aws_s3_list_objects", "params": { "bucket": "" },
        { "step": 6, "server": "python_execution", "tool": "python_execution", "params": {} },
        { "step": 7, "server": "python_execution", "tool": "python_execution", "params": {} },
        { "step": 8, "server": "slack", "tool": "send_slack_message", "params": { "channel": "" },
        { "step": 9, "server": "python_execution", "tool": "python_execution", "params": {} },
        { "step": 10, "server": "jira", "tool": "create_ticket", "params": { "project": "" },
        { "step": 11, "server": "slack", "tool": "send_slack_message", "params": { "channel": "" },
      ],
    },
    "analysis_rubric": {
      "steps": [
        { "step": 1, "extract": ["buckets[]"], "compute": [], "select": [], "accept_if": [] },
        { "step": 2, "extract": ["big_buckets[]"], "compute": [], "select": [], "accept_if": [] },
        { "step": 3, "extract": ["objects_json"], "compute": [], "select": [], "accept_if": [] },
        { "step": 4, "extract": ["error_pct1"], "compute": [], "select": [], "accept_if": [] },
        { "step": 5, "extract": ["objects_json2"], "compute": [], "select": [], "accept_if": [] },
        { "step": 6, "extract": ["error_pct2"], "compute": ["error_pct_map = merge_map("], "select": [], "accept_if": [] },
        { "step": 7, "extract": ["histogram_path"], "compute": [], "select": [], "accept_if": [] },
        { "step": 8, "extract": [], "compute": [], "select": [], "accept_if": [], "next": [] },
        { "step": 9, "extract": ["exceed_buckets[]"], "compute": [], "select": [], "accept_if": [] },
        { "step": 10, "extract": ["jira_id"], "compute": [], "select": [], "accept_if": [] },
      ]
    }
  }
}

```

```

    { "step": 11, "extract": [], "compute": [], "select": [], "accept_if": [], "next":
  ],
  "final_answer_requirements": {
    "format": "markdown",
    "must_include": ["histogram_path", "exceed_buckets"],
    "grounded_from": ["error_pct_map", "exceed_buckets", "jira_id"],
    "quality_criteria": ["include only buckets above threshold", "valid Jira ID if .
  }
},
"final_reference": {
  "answer_text": "Histogram posted. Buckets above 1%: logs-prod, archive-east. Jira
  "facts": { "exceed_buckets": ["logs-prod", "archive-east"], "jira_id": "OPS-1432",
  "citations": { "exceed_buckets": [9], "jira_id": [10], "histogram_path": [7] }
},
"judge_rubric": {
  "weights": { "coverage": 0.3, "grounding": 0.4, "clarity": 0.2, "safety": 0.1 },
  "schema": {
    "type": "object",
    "properties": {
      "coverage": {"type": "number", "minimum": 0, "maximum": 1},
      "grounding": {"type": "number", "minimum": 0, "maximum": 1},
      "clarity": { "type": "number", "minimum": 0, "maximum": 1},
      "safety": { "type": "number", "minimum": 0, "maximum": 1},
      "total": { "type": "number", "minimum": 0, "maximum": 1}
    },
    "required": ["coverage", "grounding", "clarity", "safety", "total"]
  },
  "target_length_range": [40, 160]
}
}
}
}

```

## 5) How this plugs into SkyRL

- **Dataset:** compact seed prompt + rich `ground_truth` metadata — exactly how SkyRL expects you to provide data that the env will use for rewards. ([skyrll.readthedocs.io](https://skyrll.readthedocs.io))

- **Environment:** subclass `BaseTextEnv` , parse tool calls & final answer, call your `ToolGroups`, compute rewards per turn. This is the documented pattern for multi-turn tasks with tools and final text (see **Search** example: “agent has n turns, outputs `<answer>` at the end, receive reward”). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
  - **Training:** run GRPO in SkyRL-Train exactly like their examples; the final answer reward is just another term. (If you’re using vLLM for rollout, ensure you capture per-token logprobs and refresh LoRA adapters after updates.)
- 

## 6) Gotchas & best practices

- **Leakage:** never put `final_reference` in observations; it’s only for reward.
  - **Cost:** cache LAJ calls by `(task_id, hash(policy_answer))` .
  - **Stability:** start with heavier heuristic weight (e.g., 0.7) and add LAJ (0.3).
  - **Style mismatch:** LAJ should judge **facts & grounding** more than surface style; include `facts` in the rubric context.
  - **Determinism:** keep DSL deterministic; if a step fails `accept_if` , penalize and keep going to preserve learning signal.
  - **Search-style tags (optional):** if you adopt `<answer>...</answer>` , add stop strings so sampler cuts generations cleanly (the Search example does this with `</answer>` ). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- 

If you paste your current `generate_with_llm.py` , I’ll return a **line-exact patch** adding:

- the extended schema,
- the MCP execution harness,
- the DSL evaluator,
- and the `final_reference` + `judge_rubric` writing logic.

That will give you a **self-contained** synthetic dataset where **plans are executed, analysis is applied, and reference answers are produced**, ready for **LAJ-based** final scoring during SkyRL training.

# Implementation plan for the data generator and environment design

Below is a **surgical, code-level plan** to (A) upgrade your

`src/dataset/llm/generate_with_llm.py` so the synthetic agent **actually executes its tool plan** and emits a **ground-truth final answer**, and (B) implement a **multi-turn SkyRL environment** that rewards (1) correct tool use + analysis each turn **and** (2) the **final text answer** using **LLM-as-a-Judge (LAJ)** (optionally combined with heuristics).

Where I reference SkyRL behavior, I cite the official docs/examples so your team can confirm alignment:

- **Dataset format** (what each record must contain). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- **BaseTextEnv** interface and how multi-turn envs are implemented. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- **Tools integration** (ToolGroup, parsing actions, calling tools). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- **Multi-turn final answers** (Search example uses `<answer>...</answer>` and stop strings). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- **LLM-as-a-Judge** (reference example + config pattern). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

I reviewed the three files you attached: `generate_with_llm.py` , `mini_agent_trajectories.py` , and `common.py` . They appear abbreviated with ellipses ( `...` ). I'll give **drop-in code blocks** and **exact insertion points** so you can merge them even if lines don't match 1:1.

---

## A) Upgrade `generate_with_llm.py` to produce *executed, grounded* final answers

**Goal:** after your LLM proposes a multi-turn tool plan, **actually run** that plan against your MCP servers (or the same local tool shims you'll use in the env), apply an **analysis DSL** (extract/compute/select/accept\_if), then **compose a reference final answer** from the derived state. The resulting dataset entries add:

- `ground_truth.tool_sequence` (what to do),

- `ground_truth.analysis_rubric` (how to check each step),
- `ground_truth.final_reference` (answer text + facts + citations),
- `ground_truth.judge_rubric` (weights + JSON schema the env passes to the Judge).


SkyRL's dataset loader wants **compact prompts** + **reward\_spec** per sample; the rest is your metadata for the environment to compute rewards. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

## A.1 Minimal schema extension

Add two blocks to your task schema + final writer:

- `final_answer_requirements` (format, must\_include, grounded\_from, criteria)
- `judge_rubric` (weights + structured output schema + optional length range)

These are used for **data generation** (to force the planner to think about the end state) **and** later by the **environment** to score final answers.

 **Patch: add/replace these pieces in**  
**`src/dataset/llm/generate_with_llm.py`**

### 1) Imports & helpers (top of file, after existing imports)

```
# === NEW: imports for executing plan & composing final ===
from dataclasses import dataclass
from copy import deepcopy

# If you already have a ToolManager for MCP in your repo, import it here.
# Otherwise, you can stub `execute_tool_fqn(tool_fqn: str, params: dict) -> dict`
# and later wire it to your MCP client used by the environment.
try:
    from src.utils.tool_manager import ToolManager
    MCP_AVAILABLE = True
except Exception:
    MCP_AVAILABLE = False
    ToolManager = None
```

### 2) Extend your TASK\_SCHEMA (replace/augment your dict)

```

TASK_SCHEMA = {
    "name": "skyr_task",
    "schema": {
        "type": "object",
        "properties": {
            "task_id": {"type": "string"},
            "user_prompt": {"type": "string"},
            "complexity": {"type": "string", "enum": ["simple", "moderate", "complex"]},
            "max_turns": {"type": "integer", "minimum": 2, "maximum": 20},
            "tools_available": {"type": "array", "items": {"type": "string"}},
            "limits": {"type": "object"},
            "tool_sequence": {
                "type": "array",
                "minItems": 2,
                "maxItems": 16,
                "items": {
                    "type": "object",
                    "properties": {
                        "step": {"type": "integer", "minimum": 1},
                        "server": {"type": "string"},
                        "tool": {"type": "string"},
                        "params": {"type": "object"},
                        "analysis_requirements": {
                            "type": "object",
                            "properties": {
                                "extract": {"type": "array", "items": {"type": "string"}},
                                "compute": {"type": "array", "items": {"type": "string"}},
                                "select": {"type": "array", "items": {"type": "string"}},
                                "accept_if": {"type": "array", "items": {"type": "string"}},
                                "next_args_from": {"type": "string"}
                            }
                        },
                    },
                    "required": ["next_args_from"]
                },
            },
            "required": ["step", "server", "tool", "params", "analysis_requirements"],
            "additionalProperties": True
        },
    },
    "final_answer_requirements": {

```

```

        "type": "object",
        "properties": {
            "format": {"type": "string"}, # "text" | "markdown" | "json"
            "must_include": {"type": "array", "items": {"type": "string"}},
            "grounded_from": {"type": "array", "items": {"type": "string"}},
            "quality_criteria": {"type": "array", "items": {"type": "string"}}
        },
        "required": ["format", "must_include", "grounded_from"]
    },
    "judge_rubric": {
        "type": "object",
        "properties": {
            "weights": {"type": "object"},
            "schema": {"type": "object"}, # JSON schema for LAJ structure
            "target_length_range": {"type": "array", "items": {"type": "integer"}},
        },
        "required": ["weights", "schema"]
    }
},
"required": ["task_id", "user_prompt", "complexity", "max_turns", "tool_sequence",
             "final_answer_requirements", "judge_rubric"],
"additionalProperties": True
}
}

```

**3) Add an execution record and DSL to evaluate steps** (near bottom or a new section)



```

# === NEW: Executed plan outputs ===
@dataclass
class ExecStep:
    step: int
    tool_fqn: str
    args: dict
    result_summary: dict
    accept_pass: bool
    checks: dict

@dataclass
class ExecOut:
    state: dict
    steps: List[ExecStep]

# --- Safe DSL utilities (expand as needed; keep deterministic) ---
def _extract_path(result: Any, path: str) -> Tuple[Optional[Any], bool]:
    """
    Supports simple paths like 'field', 'field[]', 'obj[][title]', 'map{key->val}' summary
    Return (value, ok)
    """
    try:
        if path.endswith("[]"): # list extraction
            key = path[:-2]
            return result.get(key, []), True
        if path.endswith("[][title]"):
            key = path.split("[]")[0]
            items = result.get(key, [])
            return [it.get("title") for it in items if isinstance(it, dict)], True
        if "{title->score}" in path:
            # summarized maps from list of items having title/score
            base = path.split("{")[0]
            items = result.get(base, [])
            return {it["title"]: it.get("score", 0.0) for it in items if "title" in it}, True
        return result.get(path, None), (path in result)
    except Exception:
        return None, False

def _compute(expr: str, state: dict) -> dict:

```

```

"""
Tiny, whitelisted DSL. Examples:
- "pct = pct_change_last_day(price_json)"
- "top3 = topk(pct, 3)"
- "tickers = regex_extract_all('[A-Z]{1,5}', content)"
- "neg_titles = head(neg_titles, 5)"
"""

out = {}
name, rhs = [s.strip() for s in expr.split("=", 1)]
def pct_change_last_day(price_json):
    # price_json: {ticker: [ {open,close,...}, ... ]}
    # return {ticker: pct} for last 2 rows
    pct = {}
    for k, arr in price_json.items():
        if len(arr) >= 2 and "close" in arr[-1] and "close" in arr[-2]:
            b, a = float(arr[-2]["close"]), float(arr[-1]["close"])
            if b != 0:
                pct[k] = (a / b) - 1.0
    return pct
def topk(d: dict, k: int):
    return [k_ for k_, _ in sorted(d.items(), key=lambda kv: kv[1], reverse=True)[:k]]
def head(lst: list, n: int): return lst[:n]
def unique(lst: list): return list(dict.fromkeys(lst))
def concat(*lsts): out=[]; [out.extend(_l) for _l in lsts]; return out
def count_keys(d: dict): return len(d.keys()) if isinstance(d, dict) else 0
def regex_extract_all(pattern: str, text: str):
    import re
    return re.findall(pattern, text or "")
# ---- Eval RHS in controlled namespace ----
safe_ns = {
    **deepcopy(state),
    "pct_change_last_day": pct_change_last_day,
    "topk": topk, "head": head, "unique": unique, "concat": concat,
    "count_keys": count_keys, "regex_extract_all": regex_extract_all,
}
# Support calls in the form fn(state_key, ...), where state_key was introduced earlier
value = eval(rhs, {"__builtins__": {}}, safe_ns) # guarded namespace
out[name] = value
return out

```

```

def _check(cond: str, state: dict) -> bool:
    # e.g., "len(tickers) >= 80", "top_gainer in tickers", "histogram_path ~= '/tmp/.*\\"
    try:
        if " ~= " in cond:
            lhs, pattern = [s.strip() for s in cond.split(" ~= ", 1)]
            import re
            return re.search(pattern.strip("'\""), str(eval(lhs, {"__builtins__": {}}, state)))
        return bool(eval(cond, {"__builtins__": {}}, state))
    except Exception:
        return False

def _resolve_placeholders(obj: Any, state: dict) -> Any:
    # Replace ${var} or ${arr[i]} inside params
    if isinstance(obj, str):
        import re
        def repl(m):
            key = m.group(1)
            try:
                return str(eval(key, {"__builtins__": {}}, state))
            except Exception:
                return m.group(0)
        return re.sub(r"\$\{([^\}]+\)}", repl, obj)
    if isinstance(obj, dict):
        return {k: _resolve_placeholders(v, state) for k, v in obj.items()}
    if isinstance(obj, list):
        return [_resolve_placeholders(v, state) for v in obj]
    return obj

```

#### 4) Execute the plan + build a final reference

```

async def simulate_plan_and_collect(task: dict, tm: Optional[ToolManager]) -> ExecOut:
    """
    Runs tool_sequence with placeholder resolution and applies analysis_requirements
    to produce named state + per-step summaries.
    """
    state: dict = {}
    exec_steps: List[ExecStep] = []
    for step_obj in task["tool_sequence"]:
        step = int(step_obj["step"])
        tool_fqn = f'{step_obj["server"]}.{step_obj["tool"]}'
        params = _resolve_placeholders(step_obj.get("params", {}), state)

        if tm is None:
            # If MCP ToolManager not wired yet, just mock a stable shape
            result = {"ok": True, "echo": params}
        else:
            result = await tm.execute_tool(tool_fqn, params, timeout=20.0)

        # Apply analysis requirements
        ar = step_obj.get("analysis_requirements", {})
        updates = {}
        missing = []
        accept = True
        for need in ar.get("extract", []):
            val, ok = _extract_path(result, need)
            if ok:
                key = need.split("[")[0].split("{")[0]
                updates[key] = val
            else:
                missing.append(need)
                accept = False
        for expr in ar.get("compute", []):
            try: updates.update(_compute(expr, {**state, **updates}))
            except Exception: accept = False
        for expr in ar.get("select", []):
            try: updates.update(_compute(expr, {**state, **updates}))
            except Exception: accept = False
        for cond in ar.get("accept_if", []):
            if not _check(cond, {**state, **updates}):

```

```

        accept = False

    state.update(updates)
    exec_steps.append(ExecStep(
        step=step, tool_fqn=tool_fqn, args=params,
        result_summary={"keys": list(result)[:10]}, # keep tiny
        accept_pass=accept,
        checks={"missing": missing, "updated": list(updates.keys())}
    ))
    return ExecOut(state=state, steps=exec_steps)

async def compose_reference_answer(task: dict, exec_out: ExecOut, client: AsyncOpenAI) ->
    far = task["final_answer_requirements"]
    facts = {name: exec_out.state.get(name) for name in far.get("grounded_from", [])}
    # Compose with a small LLM, STRICT JSON output to reduce drift
    system = "You are a concise analyst. Write the final answer strictly grounded in the p
    user = {
        "facts": facts,
        "must_include": far.get("must_include", []),
        "format": far.get("format", "text"),
        "quality_criteria": far.get("quality_criteria", []),
    }
    # If you prefer determinism or cost-free, you can template instead of calling the LLM
    resp = await client.chat.completions.create(
        model=os.getenv("OPENAI_RESP_MODEL", "gpt-4o-mini"),
        response_format={"type": "json_object"},
        messages=[
            {"role": "system", "content": system},
            {"role": "user", "content": json.dumps(user)}
        ],
        temperature=0.1,
    )
    data = json.loads(resp.choices[0].message.content)
    answer_text = data.get("answer") or data.get("final_answer") or json.dumps(facts)
    # Build simple citations: map each fact name to the last step that updated it
    name_to_step = {}
    for s in reversed(exec_out.steps):
        for k in exec_out.state.keys():
            if k not in name_to_step and k in s.checks.get("updated", []):
                name_to_step[k] = s.step

```

```
citations = {k: [name_to_step.get(k)] for k in facts.keys() if name_to_step.get(k)}  
return {"answer_text": answer_text, "facts": facts, "citations": citations}
```

**5) Thread it into your existing `_one_task(...)` flow** (after you get the *planned* task from the LLM)

```
async def _one_task(...):  
    task = await _call_llm_with_schema(..., TASK_SCHEMA, ...) # your existing planner call  
    task = _verify_and_repair(task) # ensure next_args_from chains exist (implement simple repair)  
  
    tm = ToolManager() if MCP_AVAILABLE else None # wire to MCP servers if available  
    exec_out = await simulate_plan_and_collect(task, tm)  
  
    openai_client = AsyncOpenAI(api_key=os.environ["OPENAI_API_KEY"])  
    final_ref = await compose_reference_answer(task, exec_out, openai_client)  
  
    task["_exec_out"] = {"steps": [asdict(s) for s in exec_out.steps]} # optional, keep track of execution  
    task["_final_reference"] = final_ref  
    return task
```

**6) Write the SkyRL-compatible record** (modify your writer)

```

def to_skyrl_sample(task: dict, system_prompt: str, env_class: str, data_source: str) -> dict:
    prompt_messages = [
        {"role": "system", "content": system_prompt},
        {"role": "user", "content": task["user_prompt"]},
    ]
    ground_truth = {
        "task_id": task["task_id"],
        "complexity": task["complexity"],
        "max_turns": task["max_turns"],
        "limits": task.get("limits", {}),
        "tool_sequence": task["tool_sequence"],
        "analysis_rubric": {
            "steps": [
                {"step": s["step"], **s["analysis_requirements"]}
                for s in task["tool_sequence"]
            ],
            "final_answer_requirements": task["final_answer_requirements"]
        },
        "final_reference": task["_final_reference"],
        "judge_rubric": task["judge_rubric"],
    }
    return {
        "data_source": data_source,
        "env_class": env_class,
        "prompt": prompt_messages,
        "reward_spec": {"method": "rule", "ground_truth": ground_truth},
        "extra_info": {"gen": {"model": task.get("_model"), "backend": task.get("_backend")}}
    }

```

That's all you need on the data-gen side. Each sample is now **grounded**: the plan was **executed**, the **state** was derived by your DSL, and a **reference final answer** exists for the **Judge** to compare against during training. (Dataset remains within the SkyRL format: minimal prompt + `reward_spec` — everything else lives under `ground_truth` for env reward.) ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

---

## B) Implement the multi-turn environment with per-turn rewards + final text reward

Create `src/envs/mcp_tool_env.py` (or extend your existing one) as a subclass of `BaseTextEnv` (text-in/text-out). In each step:

### 1. Parse action:

- either a **tool call**: `{"tool": "server.tool", "arguments": {...}}`
- or a **final answer**: `{"final_answer": "...text..."}`

(You can also allow `<answer>...</answer>` or `<tool>...</tool>` tags; if you do, add **stop strings** like the Search example: `stop='["</tool>", "</answer>"]'` in generator sampling.) ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

### 2. Tool step:

- resolve placeholders from the **current named state**, call the right ToolGroup method, run **analysis rubric**:  
`extract/compute/select/accept_if/next_args_from .`
- compute **shaped reward** (+ penalties if malformed).

### 3. Final step:

- compute **Heuristic** score (coverage/grounding/clarity/safety vs the requirements + reference facts).
- compute **LAJ** score by sending a compact rubric to the Judge model with **structured outputs** (numbers-only). SkyRL already documents the judge pattern (for GSM8K) — we just do a graded version with multiple dimensions.  
([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

Here's a concise implementation skeleton:



```

# src/envs/mcp_tool_env.py
from __future__ import annotations
import json, re, os, asyncio
from typing import Any, Dict, List, Tuple, Optional
from skyrll_gym.envs.base_text_env import BaseTextEnv, BaseTextEnvStepOutput # API per doc
from skyrll_gym.tools.core import ToolGroup # or your ToolManager wrappers
from openai import AsyncOpenAI

class MCPToolEnv(BaseTextEnv):
    """Multi-turn, multi-tool env that also rewards a final text answer."""
    def __init__(self, env_config: Dict[str, Any] = None, extras: Dict[str, Any] = None):
        super().__init__()
        self.turn = 0
        self.max_turns = (env_config or {}).get("max_turns", 8)
        self.reward_weights = (env_config or {}).get("reward_weights", {
            "tool_name": 0.2, "param_binding": 0.15, "extract": 0.15,
            "compute": 0.15, "accept_if": 0.1, "penalty": -0.1,
            "final_heur": 0.6, "final_laj": 0.4
        })
        # Initialize ToolGroups (search, python, finance, slack, etc.)
        # Or bridge to your MCP servers if you already have a ToolManager.
        self.init_tool_groups(self._make_tool_groups(extras or {}))
        # Judge
        self.judge = AsyncOpenAI(api_key=os.environ.get("OPENAI_API_KEY"))

        # Set by init()
        self.gt = None # ground_truth (dict from reward_spec)
        self.state = {} # named values we derive per step

    def _make_tool_groups(self, extras: Dict[str, Any]) -> List[ToolGroup]:
        # TODO: return actual tool groups; you can wrap MCP servers as ToolGroup adapters
        return []

    # ----- BaseTextEnv API -----
    def init(self, prompt, ground_truth=None) -> Tuple[list, dict]:
        """
        Called once per episode with the dataset prompt and its ground_truth from reward_spec
        """
        self.turn = 0

```

```

self.state = {}
self.gt = ground_truth or {}
return prompt, {"task_id": self.gt.get("task_id")}

def step(self, action: str) -> BaseTextEnvStepOutput:
    self.turn += 1
    kind, payload = self._parse_action(action)
    if kind == "tool":
        tool_fqn, args = payload["name"], payload.get("arguments", {})
        # Execute the tool via BaseTextEnv helper
        try:
            group_name, tool_name = tool_fqn.split(".", 1)
            obs = self._execute_tool(group_name, tool_name, [args]) # returns serial
        except Exception as e:
            # malformed tool -> penalty, continue
            return BaseTextEnvStepOutput(
                observations=[{"role": "user", "content": json.dumps({"error": str(e)})},
                reward=self.reward_weights["penalty"],
                done=(self.turn >= self.max_turns),
                metadata={"error": "tool_exec", "exception": str(e)}
            )
        # Apply analysis rubric for this step (if present)
        step_idx = self._match_step(tool_fqn)
        r_step, meta = self._score_tool_step(step_idx, tool_fqn, args, obs)
        done = (self.turn >= self.max_turns)
        return BaseTextEnvStepOutput(
            observations=[{"role": "user", "content": json.dumps(obs)[:2048]}],
            reward=r_step,
            done=done,
            metadata={"step": step_idx, **meta}
        )

    # Final answer branch
    final_text = payload
    r_final, meta = asyncio.get_event_loop().run_until_complete(
        self._score_final(final_text)
    )
    return BaseTextEnvStepOutput(
        observations=[],
        reward=r_final,

```

```

        done=True,
        metadata={"final": meta}
    )

# ----- Parsing -----
def _parse_action(self, text: str) -> Tuple[str, Any]:
    # JSON first
    try:
        obj = json.loads(text)
        if "tool" in obj:
            return "tool", {"name": obj["tool"], "arguments": obj.get("arguments", {})}
        if "final_answer" in obj:
            return "final", obj["final_answer"]
    except Exception:
        pass

    # Tag-style, compatible with Search pattern
    if "<answer>" in text and "</answer>" in text:
        ans = text.split("<answer>")[1].split("</answer>")[0].strip()
        return "final", ans

    if "<tool>" in text and "</tool>" in text:
        inner = text.split("<tool>")[1].split("</tool>")[0]
        m = re.search(r"<([a-zA-Z0-9_.-]+)>(.*?)</\1>", inner, re.DOTALL)
        if m:
            name, args_text = m.group(1), m.group(2).strip()
            try:
                args = json.loads(args_text)
            except Exception:
                args = {"raw": args_text}
            return "tool", {"name": name, "arguments": args}

    # Default: treat as final free text
    return "final", text.strip()

# ----- Rewarding -----
def _match_step(self, tool_fqn: str) -> int:
    for s in self.gt.get("tool_sequence", []):
        if f'{s["server"]}.{s["tool"]}' == tool_fqn:
            return int(s["step"])
    return self.turn # fallback

def _score_tool_step(self, step_idx: int, tool_fqn: str, args: dict, result: dict) ->

```

```

ar = None
for s in self.gt.get("analysis_rubric", {}).get("steps", []):
    if int(s["step"]) == step_idx:
        ar = s; break
if ar is None:
    # No rubric -> small neutral reward
    return 0.0, {"warn": "no_rubric"}

reward = 0.0
meta = {"tool": tool_fqn, "args": args, "accept_if": []}

# Tool choice reward
expected_fqn = None
for st in self.gt["tool_sequence"]:
    if int(st["step"]) == step_idx:
        expected_fqn = f'{st["server"]}.{st["tool"]}'; break
if expected_fqn == tool_fqn:
    reward += self.reward_weights["tool_name"]

# Param binding reward (did args use the prior state's 'next_args_from' value?)
naf = ar.get("next_args_from")
if naf:
    used = json.dumps(args)
    if naf in used:
        reward += self.reward_weights["param_binding"]

# Extract
ext_ok = True
for need in ar.get("extract", []):
    val = self._extract_path(result, need)[0]
    if val is None: ext_ok = False
    else:
        self.state[need.split("[")[0].split("{")[0]] = val
if ext_ok: reward += self.reward_weights["extract"]

# Compute/select
try:
    for expr in ar.get("compute", []): self.state.update(self._compute(expr))
    for expr in ar.get("select", []): self.state.update(self._compute(expr))
    reward += self.reward_weights["compute"]

```

```

except Exception:
    pass

# accept_if
all_ok = True
for cond in ar.get("accept_if", []):
    ok = self._check(cond)
    meta["accept_if"].append({"cond": cond, "ok": ok})
    all_ok = all_ok and ok
if all_ok: reward += self.reward_weights["accept_if"]

return reward, meta

async def _score_final(self, text: str) -> Tuple[float, dict]:
    # Heuristic
    h_score, h_meta = self._score_final_heur(text)
    # LAJ
    j_score, j_meta = await self._score_final_laj(text)
    w = self.reward_weights
    total = w["final_heur"] * h_score + w["final_laj"] * j_score
    return float(total), {"heur": h_meta, "laj": j_meta}

def _score_final_heur(self, text: str) -> Tuple[float, dict]:
    far = self.gt["analysis_rubric"]["final_answer_requirements"]
    jr = self.gt["judge_rubric"]
    weights = jr["weights"]
    lo, hi = (jr.get("target_length_range") or [0, 10**9])

    # Coverage: all must_include items appear or can be verified from state
    must = far.get("must_include", [])
    cov_hits = sum(1 for k in must if k in text or k in json.dumps(self.state))
    coverage = cov_hits / max(1, len(must))

    # Grounding: mentions consistent with reference facts
    facts = self.gt["final_reference"]["facts"]
    grounding = self._grounding_score(text, facts)

    # Clarity: simple length band
    words = len(text.split())
    clarity = 1.0 if lo <= words <= hi else 0.5 if (0.7*lo) <= words <= (1.5*hi) else

```

```

# Safety: simple heuristic
safety = 1.0 if not re.search(r"\b(SSN|password|api_key)\b", text, re.I) else 0.0

total = (weights.get("coverage",0)*coverage +
         weights.get("grounding",0)*grounding +
         weights.get("clarity",0)*clarity +
         weights.get("safety",0)*safety)
return float(total), {"coverage":coverage, "grounding":grounding, "clarity":clarity}

async def _score_final_laj(self, text: str) -> Tuple[float, dict]:
    # Compact state summary + reference
    facts = self.gt["final_reference"]["facts"]
    ref = self.gt["final_reference"]["answer_text"]
    schema = self.gt["judge_rubric"]["schema"]
    rubric_prompt = {
        "instructions": [
            "Score the FINAL answer vs the reference & facts.",
            "Return ONLY JSON with fields in the provided schema."
        ],
        "facts": facts, "reference": ref, "final": text
    }
    resp = await self.judge.chat.completions.create(
        model=os.getenv("OPENAI_JUDGE_MODEL", "gpt-4o-mini"),
        response_format={"type": "json_schema", "json_schema": {"name": "judge_schema"}},
        messages=[
            {"role":"system","content":"You are a strict evaluator."},
            {"role":"user","content": json.dumps(rubric_prompt)}
        ],
        temperature=0.0,
    )
    data = json.loads(resp.choices[0].message.content)
    return float(data.get("total", 0.0)), data

# --- local copies of DSL helpers (same semantics as generator) ---
def _extract_path(self, res: dict, path: str): return _extract_path(res, path)
def _compute(self, expr: str): return _compute(expr, self.state)
def _check(self, cond: str): return _check(cond, self.state)
def _grounding_score(self, text: str, facts: dict) -> float:
    # Example: enforce that any ticker in text belongs to facts['top3'] if that exists

```

```

import re
if "top3" in facts:
    mentions = set(re.findall(r"\b[A-Z]{1,5}\b", text))
    target = set(facts["top3"])
    if not mentions: return 0.5 # neutral
    ok = all(m in target for m in mentions if m.isupper())
    return 1.0 if ok else 0.0
return 0.5

```

- This adheres to the **BaseTextEnv** interface (returning `BaseTextEnvStepOutput` ) and integrates **ToolGroups** the same way the docs show. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- If you choose to use tags ( `<answer>` ), add `generator.sampling_params.stop=['</tool>', "</answer>"]` in your training configs, same as the **Search** example. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- If you prefer **JSON-only** actions, you can skip stop strings.

---

## C) Sample dataset items (now with `final_reference` + `judge_rubric` )

Put these in `data/processed/train_llm.json` along with your existing records; they use the same shape your environment expects:

```

{
  "data_source": "synthetic/llm",
  "env_class": "MCPToolEnv",
  "prompt": [
    { "role": "system", "content": "You have tools DuckDuckGo, yahoo_finance, python_execution" },
    { "role": "user", "content": "Find top-3 NASDAQ-100 gainers today, collect 5 headlines" }
  ],
  "reward_spec": {
    "method": "rule",
    "ground_truth": {
      "task_id": "nasdaq100_neg_digest_v2",
      "complexity": "complex",
      "max_turns": 12,
      "limits": { "max_servers": 3, "max_tools": 10 },
      "tool_sequence": [
        { "step": 1, "server": "DuckDuckGo", "tool": "search", "params": { "query": "NASDAQ" } },
        { "step": 2, "server": "DuckDuckGo", "tool": "fetch_content", "params": { "url": "https://www.duckduckgo.com" } },
        { "step": 3, "server": "yahoo_finance", "tool": "get_yfinance_price_history", "params": { "symbol": "NVDA" } },
        { "step": 4, "server": "python_execution", "tool": "python_execution", "params": { "code": "import sys; print(sys.argv[1])" } },
        { "step": 5, "server": "DuckDuckGo", "tool": "search", "params": { "query": "${top3}" } },
        { "step": 6, "server": "DuckDuckGo", "tool": "search", "params": { "query": "${top3}" } },
        { "step": 7, "server": "DuckDuckGo", "tool": "search", "params": { "query": "${top3}" } },
        { "step": 8, "server": "python_execution", "tool": "python_execution", "params": { "code": "import sys; print(sys.argv[1])" } },
        { "step": 9, "server": "python_execution", "tool": "python_execution", "params": { "code": "import sys; print(sys.argv[1])" } },
        { "step": 10, "server": "slack", "tool": "send_slack_message", "params": { "channel": "#channel" } }
      ],
    },
    "analysis_rubric": {
      "steps": [],
      "final_answer_requirements": {
        "format": "markdown",
        "must_include": ["top3", "neg_titles"],
        "grounded_from": ["top3", "title_sentiment_map"],
        "quality_criteria": ["relevant headlines", "no hallucinated tickers", "concise & clear"]
      }
    },
    "final_reference": {
      "answer_text": "Top-3: NVDA, AMD, META. Negative headlines: '...', '...'.",
      "facts": { "top3": ["NVDA", "AMD", "META"], "neg_titles": ["...", "..."] },
      "citations": { "top3": [4], "neg_titles": [8, 9] }
    }
  }
}

```



```

    },
    "judge_rubric": {
      "weights": { "coverage": 0.35, "grounding": 0.4, "clarity": 0.15, "safety": 0.10 },
      "target_length_range": [40, 140],
      "schema": {
        "type": "object",
        "properties": {
          "coverage": {"type": "number", "minimum": 0, "maximum": 1},
          "grounding": {"type": "number", "minimum": 0, "maximum": 1},
          "clarity": {"type": "number", "minimum": 0, "maximum": 1},
          "safety": {"type": "number", "minimum": 0, "maximum": 1},
          "total": {"type": "number", "minimum": 0, "maximum": 1}
        },
        "required": ["coverage", "grounding", "clarity", "safety", "total"]
      }
    }
  }
}

```

This conforms to SkyRL's dataset contract (compact prompt + env\_class + reward\_spec), leaving your **long-horizon supervision** under `ground_truth`. ([skyrllib.readthedocs.io](https://skyrllib.readthedocs.io))

## D) Training knobs to make multi-turn + final answer work smoothly

- If you adopt **tag-style** final answers or tool calls (e.g., `<answer>...</answer>`), set **stop strings** in generator sampling the same way the **Search** example does: `generator.sampling_params.stop=['</tool>', '</answer>']` (or `"</search>" / "</answer>"`). ([skyrllib.readthedocs.io](https://skyrllib.readthedocs.io))
- Keep `generator.async_engine=true` and `generator.batched=false` for **multi-turn** async rollouts (again per Search). ([skyrllib.readthedocs.io](https://skyrllib.readthedocs.io))
- Judge model and policy model can be separate; see **LLM-as-a-Judge** example for using an OpenAI model for the judge while training a Qwen policy. ([skyrllib.readthedocs.io](https://skyrllib.readthedocs.io))

## E) Why this is mathematically sound (GRPO/PPO)

You're maximizing expected return  $\mathbb{E}_{\pi}[\sum_t r_t]$  where:

- $(r_1, \dots, r_K) = \text{shaped instrumented rewards}$  (tool choice, param binding, extraction success, computations, accept\_if), and
- $(r_{K+1}) = \lambda_h \cdot r^{\text{heur}}(y_T) + \lambda_j \cdot r^{\text{LAJ}}(y_T, y^{\wedge})$  is the **terminal reward** for the final text, with  $(y_T)$  the policy's output and  $(y^{\wedge})$  the reference.

LAJ returns a **numeric** score (0–1), so it slots directly into the return — this is exactly what the **SkyRL LLM-as-a-Judge** example does (for GSM8K it's binary, we generalize to graded). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

---

## F) Sanity tests you can add right now

- **Unit test:** pass a fake step with `extract=["content"]` → ensure `self.state["content"]` updated and reward includes `extract` term.
  - **Final answer test:** with `facts = {"top3": ["NVDA"]}`, final text `"Top-3: NVDA"` → `coverage=1`, `grounding=1` → reward close to weights sum.
  - **Judge smoke:** mock the OpenAI response to return `{"coverage":1,"grounding":1,"clarity":1,"safety":1,"total":1}` and ensure reward path works offline.
- 

## G) Where this maps to the SkyRL documentation

- **Create env** with **BaseTextEnv**, define `step` that parses tool calls and final answers → “Creating a New Environment” tutorial (and the `BaseTextEnv` API excerpt). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- **Use ToolGroups** and `_execute_tool` → “Using Tools in SkyRL-Gym” (core concepts, ToolGroup, env integration, action parsing). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- **Multi-turn with final answer** and **stop strings** → “Multi-Turn RL for Search” example

and run config. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

- **Judge integration** → “LLM as a Judge for GSM8K” example (structure, env calling judge model). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
  - **Dataset contract** → “Dataset Preparation” (data\_source, prompt, env\_class, reward\_spec). ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))
- 

## TL;DR implementation checklist

### 1. **Generator** ( `generate_with_llm.py` )

- ☒ Extend schema with `final_answer_requirements` + `judge_rubric` .
- ☒ Execute plan → `ExecOut(state, steps)` via MCP tools.
- ☒ Compose `final_reference` with an LLM (or template) grounded in `state` .
- ☒ Write `reward_spec.ground_truth` including `final_reference` + `judge_rubric` .

### 2. **Environment** ( `mcp_tool_env.py` )

- ☒ Parse action as tool or final; call tools via ToolGroups.
- ☒ Implement **per-turn rubric scoring** and **final LAJ scoring** (with structured outputs).
- ☒ Return `BaseTextEnvStepOutput` each turn.

### 3. **Training config**

- ☒ If using tags, set `generator.sampling_params.stop` .
- ☒ Enable async multi-turn rollouts as in Search.

Once you wire the ToolGroups to your MCP servers, this will let you train **long-horizon, multi-tool research agents** whose **final textual answers** are graded both **heuristically** and by an **LLM judge**, entirely within SkyRL’s standard interfaces. ([skyrl.readthedocs.io](https://skyrl.readthedocs.io))

If you want me to turn these snippets into line-exact diffs against your current files, paste the full current contents of `generate_with_llm.py` and `mcp_tool_env.py` (without `...` ) and I’ll produce a ready-to-apply patch.